University of Oslo

FYS-STK3155

Project 1

Eirik eirik@ifi.uio.no
Felix felix@ifi.uio.no
Brage Wiseth bragewi@ifi.uio.no

September 14, 2023

Contents

Abstract

This paper describes the mixtures-of-trees model, a probabilistic model for discrete multidimensional domains. Mixtures-of-trees generalize the probabilistic trees of [1] in a different and complementary direction to that of Bayesian networks. We present efficient algorithms for learning mixtures-of-trees models in maximum likelihood and Bayesian frameworks. We also discuss additional efficiencies that can be obtained when data are "sparse," and we present data structures and algorithms that exploit such sparseness. Experimental results demonstrate the performance of the model for both density estimation and classification. We also discuss the sense in which tree-based classifiers perform an implicit form of feature selection, and demonstrate a resulting insensitivity to irrelevant attributes.

Keywords: Bayesian Networks, Mixture Models, Chow-Liu Trees

1. Introduction

Probabilistic inference has become a core technology in AI, largely due to developments in graph-theoretic methods for the representation and manipulation of complex probability distributions [3]. Whether in their guise as [2] directed graphs (Bayesian networks) or as undirected graphs (Markov random fields), *probabilistic graphical models* have a number of virtues as representations of uncertainty and as inference engines. Graphical models allow a separation between qualitative, structural aspects of uncertain knowledge and the quantitative, parametric aspects of uncertainty...

Remainder omitted in this sample. See http://www.jmlr.org/papers/ for full paper.

2. Ordinary Least Squares

2.1 Discussion on Scaling

Unscaled sample design matrix fitting one-dimensional polynomial of degree 5

```
[[1.      0.      0.      0.      0.      0.     ]
[1.      0.25    0.0625  0.01562 0.00391 0.00098]
[1.      0.5     0.25    0.125   0.0625  0.03125]
[1.      0.75    0.5625  0.42188 0.31641 0.2373 ]
[1.      1.      1.      1.      1.      1.     ]]
```

Scaled sample design matrix fitting one-dimensional polynomial of degree 5

```
[[ 0.      -1.41421 -1.0171  -0.83189 -0.728   -0.66226]
[ 0.      -0.70711 -0.84758 -0.7903  -0.71772 -0.65971]
[ 0.       0.      -0.33903 -0.49913 -0.56348 -0.58075]
[ 0.       0.70711  0.50855  0.29116  0.10488 -0.0433 ]
[ 0.       1.41421  1.69516  1.83016  1.90431  1.94603]]
```

Unscaled β after fitting two-dimensional polynomial of degree 5

```
[  0.4097    7.55124   3.79304 -32.85696 -14.83669  -8.81645  45.45889
43.33221  20.70625  -7.63623 -21.24552 -51.81866  -7.53731 -29.60175
28.57282   0.73824  18.29253  10.60883  -5.52465  16.60259 -16.13743]
```

Scaled β after fitting two-dimensional polynomial of degree 5

```
[ 2.21481  1.1093  -9.93941 -3.32124 -2.6612  13.11848  8.70306  4.14949
-2.19979 -5.79259 -9.48697 -1.28166 -5.40093  7.77825  0.19055  3.09519
1.60557 -0.83485  2.797   -4.1589 ]
```

First x_{unscaled} and x_{scaled} is not that different, the original data was close to zero-centerd and not that spread out, which means that initaially by just looking at the data scaling is not that nececary. When we add the polynomial terms we can now see that some of the entries of X_{unscaled} get really small as an exsample $0.1^5 = 0.00001$ this makes the collums of X_{unscaled} live in their own order of magnitude and scaling should be considerd to bring them back to the same ish order of magnitude. The act of not scaling results in β spanning from -50 to 48 while scaling gives a smaller span from -10 to 13 (We would imagine that keeping everything in the same order of magnitude is something that the computer likes, perhaps reducing floatingpoint error). Now this alone may not justify why we should scale this dataset, as scaled and unscaled OLS yields the same MSE. However when doing ridge regression the cost function is directly dependent on the magnitude of β_i This explains why the MSE for unscaled data is much higher than for scaled data in the ridge case. This leads us to conclude that we should scale the data, making it easier to tweek λ and giving us nicer numbers to work with.

Acknowledgments

We would like to acknowledge support for this project from the National Science Foundation (NSF grant IIS-9988642) and the Multidisciplinary Research Program of the Department of Defense (MURI N00014-00-1-0637).

Appendix A.

In this appendix we prove the following theorem from Section 6.2:
Theorem *Let u, v, w be discrete variables such that v, w do not co-occur with u (i.e., $u \neq 0 \;\Rightarrow\; v = w = 0$ in a given dataset $\mathcal{D}$). Let N_{v0}, N_{w0} be the number of data points for which $v = 0, w = 0$ respectively, and let I_{uv}, I_{uw} be the respective empirical mutual information values based on the sample $\mathcal{D}$. Then*

$$N_{v0} \;>\; N_{w0} \;\;\Rightarrow\;\; I_{uv} \;\leq\; I_{uw}$$

with equality only if u is identically 0. ■
Proof. We use the notation:

$$P_v(i) \;=\; \frac{N_v^i}{N}, \quad i \neq 0; \quad P_{v0} \;\equiv\; P_v(0) \;=\; 1 - \sum_{i \neq 0} P_v(i).$$

These values represent the (empirical) probabilities of v taking value $i \neq 0$ and 0 respectively. Entropies will be denoted by H. We aim to show that $\frac{\partial I_{uv}}{\partial P_{v0}} < 0$....

Remainder omitted in this sample. See http://www.jmlr.org/papers/ for full paper.

Bibliography

[1] C. K. Chow and C. N. Liu. Approximating discrete probability distributions with dependence trees. *IEEE Transactions on Information Theory*, IT-14(3):462–467, 1968.

[2] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. Book in preparation for MIT Press.

[3] Judea Pearl. *Probabilistic Reasoning in Intelligent Systems: Networks of Plausible Inference*. Morgan Kaufman Publishers, San Mateo, CA, 1988.